

UseOfReason

Sam Harris not getting an ought from an is

Sam Harris recently made a series of Tweets which, he claimed, showed how to get an ‘ought’ from an ‘is’. Here they are:

1. Let’s assume that there are no ought’s or should’s in this universe. There is only what *is*—the totality of actual (and possible) facts.
2. Among the myriad things that exist are conscious minds, susceptible to a vast range of actual (and possible) experiences
3. Unfortunately, many experiences suck. And they don’t just suck as a matter of cultural convention or personal bias—they really and truly suck. (If you doubt this, place your hand on a hot stove and report back.)
4. Conscious minds are natural phenomena. Consequently, if we were to learn everything there is to know about physics, chemistry, biology, psychology, economics, etc., we would know everything there is to know about making our corner of the universe suck less.
5. If we *should* to do anything in this life, we should avoid what really and truly sucks. (If you consider this question-begging, consult your stove, as above.)
6. Of course, we can be confused or mistaken about experience. Something can suck for a while, only to reveal new experiences which don’t suck at all. On these occasions we say, “At first that sucked, but it was worth it!”
7. We can also be selfish and shortsighted. Many solutions to our problems are zero-sum (my gain will be your loss). But *better* solutions aren’t. (By what measure of “better”? Fewer things suck.)
8. So what is morality? What *ought* sentient beings like ourselves do? Understand how the world works (facts), so that we can avoid what sucks (values).

The whole thing boils down to premise 5. He says that we ought avoid things that ‘suck’. By ‘suck’ he basically means things that are painful (as his example of the stove indicates). So premise 5 basically just says: we ought to avoid pain. That is assuming an ought coming from an is: we ought not do things that cause pain (that ‘suck’).

The only thing he says to justify this is “If you consider this question-begging, consult your stove, as above”. But all ‘consulting the stove’ would do is remind us how painful the experience was. It wouldn’t, on its own, show us that we ‘ought’ not do it.

Advertisement



What Harris is relying on is the fact that we don’t *want* to have the experience of pain that touching the stove provides. The idea is that there is a hypothetical norm of the following form:

If you don’t want to feel pain, you ought not put your hand on the stove.

Harris is relying on the fact that we all don’t like feeling pain, and so the antecedent condition applies universally. But still, it is a hypothetical norm, not an unconditional (or ‘categorical’) norm.

What difference does that make?

Well, it isn’t really an example of getting an ought from an is; at least, not in any morally significant sense anyway. That’s because hypothetical norms are just the best ways of realising your desires. If you desire x, you *ought* to do y, when y is the optimal way of realising x. They can be morally significant things, like if you want to make the world a better place, you ought to give to charity, etc. But they can also be morally neutral: if you want to get your car fixed, you ought to take it to a mechanic; if you want to lose weight, you ought to take more exercise. They can also be immoral: if you want to murder your neighbour, you ought hit him over the head with this rock.

Morality, on the other hand, is usually thought of as being unconditional, or ‘categorical’. Take my last example. Sure, hitting your neighbour is an efficient way of murdering him. But we generally think that we simply ought not murder people at all.

Even if I want to, I ought not do it. The ‘is-ought’ issue is about how to derive these sorts of ‘oughts’ from mere ‘is’s.

So the mismatch is that he is asserting a categorical normative statement (“we should avoid what really and truly sucks”), and he is offering only a hypothetical norm as its justification (which is that *if* you don’t want to experience things that suck, you ought not do things that will produce experiences that suck).

Hypothetical norms can’t justify categorical norms though, because the former require you to have a particular desire, whereas categorical norms are independent of what you desire; hypothetical norms only apply to you if you have a certain desire, but categorical norms apply to you regardless of whether you do.

Advertisement

Its a bit like saying ‘Everything is A’, but justifying that with the statement ‘Everything which is B is A’. Even if we agree with the latter, that cannot justify believing the former.

Advertisement

 July 23, 2018  apmalpass

15 thoughts on “Sam Harris not getting an ought from an is”

 **jtveg**

July 23, 2018 at 1:01 pm

Reblogged this on [jtveg's Blog](#).

 Like

 **Siggy**

July 23, 2018 at 2:47 pm

Although I agree that premise 5 is the crux of Sam's argument, premise 7 also raises several questions. Why is selfishness bad? If our morality rests on "consulting the stove", I observe that it's only painful when I touch a hot stove, and not when you do. And where do we get this notion of "zero-sum"? The stove does not obviously provide a way to sum up or compare painful experiences. He's also glossing over game theory considerations, and I'm left wondering what his reasoning implies about, e.g., prisoner's dilemmas.

★ Like



Daniel Linford

July 23, 2018 at 3:27 pm

Alex —

Excellent post. I'm hardly one to defend Sam Harris — I find much of what he says either disagreeable or naive — but there may be two ways to repair his general thought and to construct a more sophisticated ethical system.

Consider the way that Shafer-Landau introduces utilitarianism in his intro ethics textbook. Shafer-Landau makes the claim that there may be things (states of affairs?) that are intrinsically undesirable. It's not that these things are undesirable because they bring about something else that's undesirable, but rather that they are simply undesirable as such. Pain is a prototypical example of something like that. And there may be reasons why agents coordinate to lessen the pain among any in their number; I think Singer may have identified some good reasons in his book *The Expanding Circle*. Singer notes that there are a variety of conditions under which two altruists will mutually fair better than two egoists; given a population in which there is a mutation resulting in at least two altruists, their genes will outcompete the egoistic genes. It's a neat idea.

A different way that one might try to save Harris's view — or something like it — is brought out by your distinction between hypothetical and categorical imperatives. When we adopt a hypothetical imperative, we adopt a maxim for our action. (Or so we may interpret Kant.) Some maxims might be such that, were everyone to adopt that maxim, everyone should expect to be better off (eg, to avoid more pains). So, since we each desire avoiding pain, and none of us (rationally) should consider others significantly different from oneself, we should all adopt those maxims. This move would remake Harris into a rule utilitarian.

In any case, presumably, what all of this demonstrates is that Harris's endeavor to do ethics without reading ethics is as wrongheaded as could be.

★ Like

 **apmalpass** •

July 23, 2018 at 4:15 pm

Isn't Shafer-Landau a non-naturalist though? If so, doesn't that mean that he accepts that you can't get an ought from an is? I thought his view is that non-natural goodness supervenes on various natural states, like pain. I like that view, but it's one that starts off by rejecting the whole point Harris was trying to argue for here.

Also, I'm not sure about Singer's current view, but I heard he has taken a realist turn partly due to the influence of Parfit, who was obviously also a non-naturalist.

★ Like

 **Hugh Jidiette**

July 23, 2018 at 4:27 pm

I think any realist thinks you can get an ought from an is, but just not without using a bridge principle like the golden rule, or whatever. Of course since the golden rule itself contains an ought, it wouldn't be a counterexample to Hume. So realists should accept Hume's logical point about the is-ought gap, while denying that it's important. David Brink has pointed out that if you want to derive an is-of-thermodynamics from an is-of-statistical-mechanics, you need a bridge principle like the Boyle-Charles bridge law. So it shouldn't bother moral realists that they need a bridge principle in ethics too.

I think Hume was skeptical about how we could know these bridge principles, so that is why the logical point about the is-ought gap is important. We can know about is's (matters of fact), and since oughts aren't logically entailed by is's, we can't know about oughts.

★ Like

 **Hugh Jidiette**

July 24, 2018 at 1:16 am

I think that the logical is-ought gap is not a problem for a naturalistic reduction. You can't derive conclusions about h₂o without premises that contain h₂o, but that doesn't mean that h₂o is not reducible (i.e. equivalent) to water. We don't need to posit that h₂o is sui generis and irreducible based on the logical point. I think a naturalist can accept the logical is-ought gap and remain unmoved in their naturalism.

 Like

 apmalpass 

July 24, 2018 at 12:18 pm

So H₂O = water is a necessary, a posteriori, synthetic identity, on the usual Kripkean view. Some people want to say that the same move can be made with respect to moral terms. Bilgrami

(https://philosophy.columbia.edu/files/philosophy/content/Self-Knowledge_and_Resentment.pdf) has an interesting take on this move, although directed at intentional states and dispositions. The idea is Fregean. So he says that if there is some proposed relation between 'good' and some naturalistic property, like 'maximising wellbeing' (etc), and it is claimed to be a synthetic identity (like H₂O and water), then we can say the following. Imagine someone denied that the good is maximising wellbeing. Such a person isn't being irrational necessarily. Rather, they lack a bit of worldly knowledge (just like if someone denied that Hesphorus was Phosphorus before the discovery of Venus, etc). "But if that is so, then to account for the fact that such a person's mind represents a completely consistent state of affairs, the terms on each side of the equation being denied will need to have a sense over and above a reference." (760). So the idea is that if being able to doubt the identity is coherent, there must be a meaning (or sense) for the terms involved. So what is the meaning of 'good', such that it can be distinct from 'maximising wellbeing', even if they are coextensive? Bilgrami says that either this sense will refer to a naturalistic property or a non-naturalistic property. Obviously, non-naturalistic is out of the question for such a naturalist, so they have to opt for a naturalistic property. But then they are asserting that the meaning of 'good' is some naturalistic property, which is not a synthetic identity claim any more, but an analytic one. And those sorts of claims are the ones Moore's arguments originally targeted.

If Bilgrami is right about that, the sort of move you outline isn't an option for a naturalist.

 Like

**Hugh Jidiette**

July 24, 2018 at 1:33 pm

I don't fully understand the argument there. Let's say the sense of Jesus (assuming proper names have senses) is supernatural, but that it refers to something natural (say, some mere human). I don't think I have to give up on naturalism because of that. I'm probably missing something.



Like

**apmalpass**

July 24, 2018 at 2:07 pm

I'll try to explain it again.

Let's say that you conceded Moore's open question argument, but you were still a determined naturalist, and you had read a bit of Kripke, and wanted to give that a try here. So you said that 'good' and 'maximising wellbeing' are not analytically identical (they don't mean the same thing), but they are coextensive, just like water and H₂O. Good = maximising wellbeing is a synthetic identity, not an analytic identity. Ok.

I can coherently doubt the identity of coextensive but analytically distinct terms; Lois Lane can doubt that Clark Kent is Superman without being irrational (she is just ignorant of some worldly fact). But to account for her being able to do this, it must be that she is using the senses of the terms, not the references. In her head she is thinking about the glasses guy and the red cape guy, and these ideas are distinct. So to be able to coherently doubt two coextensive terms, they must each have senses, and they must be distinct from one another. So because we can coherently doubt the synthetic identity of the good and maximising wellbeing, there must be a sense to 'good'. It must mean something. So the naturalist who is trying this Kripkean strategy is forced (by the coherence of denial) to admit that the terms in his equation have distinct senses, as well as coextensive references.

If I say that the sense of 'water' is clear, colorless liquid, running and rivers and streams (or whatever), then clearly I have in mind a collection or cluster of natural properties. So it could be that when the naturalist says 'good' what they mean is some cluster of natural properties like this as well. Maybe they mean helping old

ladies across roads, and giving to donkey charities or something. But if they do, then they already think that the good *means* something natural (like, that it is analytically identical to that cluster of properties). If so, then we can ask “I accept that this is helping an old lady across the road, but is it good?”, i.e. we can run Moore’s open question argument. As we set this up by saying that the naturalist accepts Moore’s open question argument, and that this is why they are trying the synthetic identity claim instead, finding themselves here is game over.

The only other option though is to say that it is a non-natural property (or cluster of). But if they think that the sense of ‘the good’ is non-natural, then they basically concede non-naturalism from the get-go. They are trying to reduce the good to natural things after all.

 Liked by 1 person

 **Hugh Jidiette**

July 24, 2018 at 2:32 pm

What do you think about my Jesus example?

I thought the synthetic naturalist accepted that good is semantically primitive from the get-go. So they wouldn’t say that the sense of good is helping an old lady across the street. The sense of good is the same as what everyone else is using. If the senses of good varied from person to person, then there couldn’t be moral disagreement, since people would be talking about two different things. The only thing a synthetic reductionist is trying to reduce is the metaphysics of it and not the semantics of it. (Other naturalists, like Cornell realists, reduce neither the semantics nor the metaphysics of it.)

 Like

 **Hugh Jidiette**

July 23, 2018 at 4:21 pm

Hypothetical norms can’t be used to refute Hume’s is-ought gap, since the hypothetical norm contains an ought within it. The point for Hume was that if an ought was in the conclusion, you would see an ought somewhere in the premises.

Charles Pigden writes:

"[Hume] was merely denying the existence of logically valid arguments from the non-moral to the moral. This becomes clear once we note that Hume does not think that he has to argue for the apparent inconceivability of is/ought deductions. It is something he thinks he can take for granted. This is what we would expect if he were making the logical point since it would have been obvious to his readers. For it was a commonplace of Eighteenth Century logical theory that in a logically valid argument the matter of the conclusion – that is the non-logical content – is contained within the premises, and thus that you cannot get out what you haven't put in. Thus if an 'ought' appears in the conclusion of an argument but not in the premises, the inference cannot be logically valid. You can't deduce an 'ought' from an 'is' by means of logic alone."

Hume should also say that you can't derive chair conclusions from premises without chairs. Since then, Prior has provided a counter example with:

1. Tea-drinking is common in England.
2. Therefore, Either tea-drinking is common in England or all New Zealanders ought to be shot.

//hypothetical norms only apply to you if you have a certain desire, but categorical norms apply to you regardless of whether you do.//

I think a lot of people understand it this way, but I'm not so sure. A hypothetical norm has the form "if you desire X, and φ -ing is the best way to achieve X, then you ought to X." A categorical norm has the form "you ought to φ ." But what if φ is "you ought to do whatever would satisfies your desires the best." Now that seems like a categorical norm that depends on your desires.

 Like



apmalpass

July 24, 2018 at 12:02 pm

I think maybe the objection should be rephrased actually. Maybe the real problem is that he is offering a way of supposedly deducing the categorical norm, "we ought not X", in the following way:

1. If you desire not to experience pain, you ought not do X
2. Everyone desires not to experience pain
3. Therefore, you ought not do X

Then the issue is just that it doesn't constitute a counterexample to Hume because premise 1 is normative. The more I think about it, the less it really matters to distinguish the hypothetical and categorical norms so much.

As for “you ought to do whatever would satisfies your desires the best,” that still seems to be a categorical norm. It applies to you whether you have desires or not. If you have desires, it is saying that you ought do whatever satisfies them best, obviously. What about if you have no desires? Then it is trivially satisfied, whatever you do; it can't be violated if you have no desires. Think about the obligation not to murder other people. If there are no other people, the obligation still applies to you, just its trivially satisfied by whatever you do. So I'm not seeing that as a problem example for the distinction.

★ Like



Hugh Jidiette

July 24, 2018 at 2:23 pm

Plausibly, oughts only apply to agents (as opposed to rock) as part of the concept of it. I'm not sure it makes sense to think of an agent without desires. Maybe something like Data from Star Trek or Vulcans?

If we accept the principle that ought implies can, then I don't think “do whatever satisfies your desires the best” could apply to someone without desires. It's like “you ought to run to the store with your 6 legs.” I don't have 6 legs, so I don't think that ought could apply to me. I think the norm is inapplicable rather than being trivially satisfied in the case that the agent has no desires. When the norm is applicable, then the categorical norm depends directly on what you desire. That's my current thinking, anyway.

★ Like



realisticnihilist

July 23, 2018 at 10:13 pm

What's kind boggling is that this guy got a fucking PhD. Are they just handing these things out like candy in Halloween or what? I should go pick mine up. Would look good on my resume.

★ Like



July 24, 2018 at 5:28 am

If you got a PhD in Russian literature, you could then write books about neuroscience. That seems to be how it works.

★ Liked by 2 people

Blog at WordPress.com. Do Not Sell or Share My Personal Information